Onkar Jadhav

AI & Automation Engineer
Pune, Maharashtra
B.Tech Computer Science, PCCOE (9.08 CGPA)

+91-7058604679
onkarj012@gmail.com
github.com/Onkarj012
linkedin.com/in/onkarj012

Summary

Final-year CS student (9.08 CGPA) with hands-on experience building Generative AI applications, multi-agent systems, workflow automation pipelines, and chatbots. Shipped Stride to production, a multi-agent AI platform with voice input, NLU, and five specialised domain agents that automate personalised health coaching. Built Lore Engine, a full-stack RAG chatbot ingesting 6,000+ pages with streaming responses and automated evaluation harnesses. Built PromptForge, an agentic pipeline that automates iterative prompt engineering across OpenAI, Anthropic, Gemini, and open-source LLMs via REST API integration. Strong in Python, LLMs, prompt engineering, REST APIs, SQL, and Git.

Education

•Bachelor of Technology in Computer Science *Nov 2022 – Jun 2026*
Pimpri Chinchwad College of Engineering, Savitribai Phule Pune University CGPA: 9.08 (up to 7th semester)

– Relevant Coursework: Machine Learning, Data Structures & Algorithms, Database Management Systems, Operating Systems, Computer Networks, Object-Oriented Programming

Experience

•Stride (AI-First Adaptive Wellness Platform) *April 2026 – Present*
Founder & Lead Engineer Remote

– Architected and shipped a **multi-agent AI chatbot** with 5 domain agents (diet, workout, sleep, hydration, habits) that automates personalised health coaching via natural-language conversations, live with 5–10 active trial users.
– Built an end-to-end **voice-to-structured-data automation pipeline** where Deepgram Nova-2 transcription feeds into an NLU extraction layer, converts free-text input into structured JSON, and routes it to the appropriate agent, eliminating all manual data entry for users.
– Integrated **multiple REST APIs** (Deepgram, OpenAI, nutrition/fitness data APIs) with async job queues and idempotent retries, ensuring reliable automation under LLM provider failures and network timeouts.
– Designed a **behavioral memory layer** shared across agents, enabling context-aware automated responses that adapt to each user's history without manual intervention.

Engineering Projects

•Lore Engine — *RAG Chatbot Platform with Automated Knowledge Ingestion* *2026 – Present*
Next.js, FastAPI, PostgreSQL, pgvector, BM25, Celery, Redis, Cloudflare R2, OpenAI, Anthropic Claude

– Built an **automated document ingestion pipeline** (Celery workers + Redis queue) processing **6,000+ MediaWiki pages** into a hybrid retrieval system (BM25 + pgvector), enabling a chatbot to answer precise contextual questions over large knowledge bases.
– Implemented **SQL-backed metadata filtering** at retrieval time (per-chunk spoiler/access flags stored in PostgreSQL), demonstrating SQL fundamentals applied to production AI data pipelines.
– Shipped a streaming **GenAI chatbot** with inline citations and source attribution, integrating OpenAI and Anthropic Claude APIs via REST.
– Built an **automated evaluation harness** (50-question ground-truth test suite) measuring retrieval MRR and answer accuracy, which gates every deployment and removes manual QA from the release workflow.

•PromptForge — *Agentic Prompt Engineering & LLM Workflow Automation Platform* *Jan 2026 – Feb 2026*
Node.js, Express.js, BullMQ, Convex, OpenRouter, React.js

– Built a **creator-critic-evaluator agentic workflow** that automates iterative prompt improvement by scoring clarity, completeness, constraint coverage, and token efficiency across versioned revision trails without manual iteration.
– Integrated **REST API routing** across OpenAI, Anthropic, Gemini, and open-source models via OpenRouter, automating multi-model evaluation for objective quality, latency, and cost comparisons.
– Engineered a **fault-tolerant async automation backend** (BullMQ workers, idempotent retries, Convex persistence) sustaining reliable prompt regression workflows without job loss on LLM provider timeouts.

Technical Skills

AI / GenAI: LLMs, Generative AI, Conversational AI, RAG Pipelines, Prompt Engineering, AI Agents, Agentic Workflows, NLP, Embeddings, Behavioral Memory, Voice AI, LLM Evaluation, OpenRouter
LLM Providers: OpenAI GPT-4o, Anthropic Claude, Gemini, Groq, Deepgram Nova-2
Automation: Async job pipelines (BullMQ, Celery), event-driven workflows, REST API integrations, webhook-driven automation, structured data extraction from unstructured input
Languages: Python, TypeScript, JavaScript, SQL, C++, Java
Backend: FastAPI, Node.js, Express.js, REST APIs, WebSockets, Redis, JSON
Frontend: React.js, Next.js, Tailwind CSS
Databases: PostgreSQL (+ pgvector), MongoDB, Convex, SQL query design
DevOps & Tools: Git, GitHub, AWS, Cloudflare R2, async pipeline architecture