Onkar Jadhav

Frontend Developer
Pune, Maharashtra
B.Tech Computer Science, PCCOE (9.08 CGPA)

+91-7058604679
onkarj012@gmail.com
github.com/Onkarj012
linkedin.com/in/onkarj012

Summary

Final-year CS student (9.08 CGPA) building responsive React and Next.js products with AI-backed workflows. Designed multi-agent LLM orchestration pipelines with multi-provider routing and a hybrid RAG system (BM25 + pgvector, reciprocal rank fusion) gated behind a deploy-time eval harness - the same CI/CD-for-AI pattern required in production analytics platforms. Built agentic orchestration (creator–critic–evaluator pipeline, cross-model benchmarking) and shipped a live multi-agent product with 5–10 active users. Strong in Python, FastAPI, RAG, agentic AI, and AWS.

Education

•Bachelor of Technology in Computer Science *Nov 2022 – Jun 2026*
Pimpri Chinchwad College of Engineering, Savitribai Phule Pune University CGPA: 9.08 (up to 7th semester)
– Relevant Coursework: Machine Learning, Data Structures & Algorithms, Database Management Systems, Operating Systems, Computer Networks, Object-Oriented Programming

Experience

•Stride - AI-First Adaptive Wellness Platform *April 2026 – Present*
Founder & Lead Engineer Remote
– Built and shipped Stride to production as sole engineer; running live with **5–10 active trial users** driving iterative product feedback.
– Architected five specialised domain agents (diet, workout, sleep, hydration, habits) sharing a behavioral memory layer, with multi-provider LLM routing via OpenRouter across Anthropic Claude, OpenAI GPT, and Gemini - eliminating vendor lock-in across five concurrent agent types.
– Separated LLM reasoning from deterministic computation (BMR/NEAT/EAT/TEF calculator), eliminating opaque model arithmetic errors; implemented CI-style quality gates before each capability release.

Engineering Projects

•Lore Engine - *RAG Platform with Deploy-Gated Evaluation Harness* *2026 – Present*
Next.js, FastAPI, PostgreSQL, pgvector, BM25, Celery, Redis, Cloudflare R2, OpenAI, Anthropic Claude
– Architected a full-stack RAG application ingesting **6,000+ MediaWiki pages** per series into a hybrid retrieval pipeline (BM25 + pgvector via reciprocal rank fusion) - chosen over pure vector search to handle exact entity and proper-noun matching that semantic models fail on.
– Designed section-aware chunking at H2/H3 boundaries with enriched `[Series][Page][Section]` metadata and spoiler-filter tags per chunk, demonstrating fine-grained context control over large knowledge bases.
– Established a **50-question per-series ground-truth eval harness** measuring retrieval MRR and end-to-end answer accuracy, gating every deploy - the same CI/CD-for-AI pattern required in production LLM analytics systems.
– Built an async ingestion pipeline (Celery, Redis, Cloudflare R2) that fetches from the MediaWiki API, converts wikitext to markdown, and indexes ∼**5M vectors per series** in pgvector - chosen over Pinecone/Weaviate to stay in one network hop.

•PromptForge - *Agentic LLM Orchestration & Cross-Model Benchmarking Platform* *Jan 2026 – Feb 2026*
Node.js, Express.js, BullMQ, Convex, OpenRouter, React.js
– Architected a creator–critic–evaluator agentic orchestration pipeline, scoring prompt clarity, constraint coverage, and token efficiency across versioned, regression-tracked revision trails - treating prompts as testable production artifacts.
– Implemented cross-model routing across OpenAI, Anthropic, Gemini, and open-source models via OpenRouter for objective quality, latency, and cost-per-output benchmarking.
– Built a fault-tolerant async backend (BullMQ workers, idempotent retries, Convex persistence) and emitted structured prompt outputs with token budget, latency profile, and version changelog - enabling prompt deployment like code.

Technical Skills

AI / ML: LLM Orchestration, Agentic Workflows, RAG Pipelines, Deploy-Gated Evaluation, pgvector, BM25, Reciprocal Rank Fusion, Multi-Agent Systems, Prompt Engineering, Behavioral Memory, Hallucination Detection, TensorFlow, OpenRouter
LLM Providers: Anthropic Claude, OpenAI GPT, Gemini, Groq, Deepgram Nova-2
Languages: Python, TypeScript, JavaScript, SQL, C++, Java
Backend: FastAPI, Node.js, Express.js, BullMQ, Celery, Redis, REST APIs, WebSockets
Frontend: React.js, Next.js, Tailwind CSS, HTML, CSS
Databases: PostgreSQL (+ pgvector), MongoDB, Convex
Cloud & DevOps: AWS, Cloudflare R2, Git, GitHub, async job pipelines, CI-style eval gating
System Design: Multi-agent coordination, hybrid retrieval, event-driven architecture, snapshot-first serving