Onkar Jadhav

Software Engineer
Pune, Maharashtra
B.Tech Computer Science, PCCOE (9.08 CGPA)

+91-7058604679
onkarj012@gmail.com
github.com/Onkarj012
linkedin.com/in/onkarj012

Summary

Final-year Computer Science student (9.08 CGPA) with strong programming fundamentals in **Java, Python, and C++**, hands-on **Linux** development, and solid Data Structures & Algorithms. Experienced in designing and shipping **fault-tolerant backend systems** — async job queues, idempotent retries, concurrent workers, REST APIs, and SQL databases — that stay reliable under provider failures, timeouts, and load. Comfortable owning services end to end: from schema design to deployment and automated testing. Fast learner, strong communicator, seeking a Software Engineer role building high-reliability, performance-critical systems.

Education

•Bachelor of Technology in Computer Science *Nov 2022 – Jun 2026*
Pimpri Chinchwad College of Engineering, Savitribai Phule Pune University CGPA: 9.08 (up to 7th semester)

- Relevant Coursework: Data Structures & Algorithms, Operating Systems, Computer Networks, Database Management Systems, Object-Oriented Programming, Machine Learning

Experience

•Stride — Adaptive Wellness Platform *April 2026 – Present*
Founder & Lead Engineer Remote

- Designed and shipped a production backend as sole engineer — **Python** services, REST APIs, and async job processing — live with real trial users.
- Built a **fault-tolerant integration layer** across multiple external APIs with async job queues and **idempotent retries**, keeping the system reliable under upstream failures and network timeouts.
- Modeled and persisted a shared state layer across services, applying schema design and SQL fundamentals to keep data consistent under concurrent access.

Engineering Projects

•Lore Engine — *Distributed Data Ingestion & Retrieval Backend* *2026 – Present*
Python, FastAPI, PostgreSQL, Celery, Redis, Cloudflare R2, Next.js

- Built a **distributed ingestion pipeline** (Celery workers + Redis queue) processing **6,000+ documents** concurrently into a Postgres-backed retrieval store, engineered for throughput and restartable, idempotent jobs.
- Implemented **SQL-backed indexing and metadata filtering** in PostgreSQL, tuning queries and access flags at retrieval time — applied DBMS fundamentals to a real data-serving path.
- Built an **automated evaluation/regression harness** (50-question ground-truth suite) that gates every deploy, catching correctness regressions before release without manual QA.
- Developed and debugged the full stack on **Linux**, using shell tooling, process management, and log-driven debugging across workers and services.

•PromptForge — *Fault-Tolerant Async Job Processing Backend* *Jan 2026 – Feb 2026*
Node.js, Express.js, BullMQ, Redis, Convex, React.js

- Engineered a **concurrent async backend** (BullMQ workers on Redis, idempotent retries, durable persistence) sustaining reliable job execution **without loss on upstream timeouts** — the reliability model trading systems demand.
- Designed a versioned pipeline that routes requests across multiple external providers and compares **latency, cost, and correctness**, exposing objective performance data per run.
- Built REST APIs and structured error handling so failed jobs surface cleanly and reprocess safely.

Technical Skills

Languages: Java, Python, C++, SQL, JavaScript, TypeScript
CS Fundamentals: Data Structures & Algorithms, Operating Systems, Computer Networks, DBMS, Object-Oriented Programming, System Design
Linux & Tooling: Linux/Unix development, Bash, shell scripting, process & log debugging, Git, GitHub
Backend: REST APIs, FastAPI, Node.js, Express.js, async job queues (BullMQ, Celery), Redis, WebSockets, concurrency, idempotent retries
Databases: PostgreSQL, SQL query design & indexing, MongoDB
Reliability: Fault-tolerant pipelines, automated testing/regression harnesses, event-driven architecture
Cloud: AWS, Cloudflare R2