Onkar Jadhav

AI / Full-Stack Software Engineer
Pune, Maharashtra
leo4.dev

+91-7058604679
onkarj012@gmail.com
github.com/Onkarj012
linkedin.com/in/onkarj012

Summary

AI/full-stack engineer (B.Tech CS, PCCOE, 9.17 CGPA) building full-stack AI products end to end, from React/Next.js frontends to Python/Node backend services. Architected a hybrid RAG system (Next.js + FastAPI, BM25 + pgvector, reciprocal rank fusion) gated behind a deploy-time eval harness, and an agentic creator–critic–evaluator pipeline with cross-model benchmarking. Currently shipping Stride, a five-agent AI product, as Founder & Lead Engineer. Strong across React, Next.js, FastAPI, Node.js, RAG, and agentic AI system design.

Education

•Bachelor of Technology in Computer Science *Nov 2022 – Jun 2026*
Pimpri Chinchwad College of Engineering, Savitribai Phule Pune University CGPA: 9.17

– Relevant Coursework: Machine Learning, Data Structures & Algorithms, Database Management Systems, Operating Systems, Computer Networks, Object-Oriented Programming

Experience

•Stride — AI-First Adaptive Wellness Platform *April 2026 – Present*
Founder & Lead Engineer Remote

– Architected and shipped Stride to production as sole engineer, owning the full stack — React/Next.js frontend, backend services, and data layer — running live with an early private user cohort driving iterative product feedback ahead of public beta.
– Architected five specialised domain agents (diet, workout, sleep, hydration, habits) sharing a behavioral memory layer, with multi-provider LLM routing via OpenRouter across Anthropic Claude, OpenAI GPT, and Gemini — eliminating vendor lock-in across five concurrent agent types.
– Separated LLM reasoning from deterministic computation (BMR/NEAT/EAT/TEF calculator), eliminating opaque model arithmetic errors; implemented CI-style quality gates before each capability release.
– Selected Convex for real-time reactive state over a REST + WebSockets stack, keeping five concurrent agent types in sync across the frontend without manual polling logic.

Engineering Projects

•Lore Engine — *RAG Platform with Deploy-Gated Evaluation Harness* *2026 – Present*
Next.js, FastAPI, PostgreSQL, pgvector, BM25, Celery, Redis, Cloudflare R2, OpenAI, Anthropic Claude

– Architected a full-stack RAG application ingesting **6,000+ MediaWiki pages** per series into a hybrid retrieval pipeline (BM25 + pgvector via reciprocal rank fusion) — chosen over pure vector search to handle exact entity and proper-noun matching that semantic models fail on.
– Designed section-aware chunking at H2/H3 boundaries with enriched `[Series][Page][Section]` metadata and spoiler-filter tags per chunk, demonstrating fine-grained context control over large knowledge bases.
– Established a **50-question per-series ground-truth eval harness** measuring retrieval MRR and end-to-end answer accuracy, gating every deploy — the same CI/CD-for-AI pattern required in production LLM analytics systems.
– Built an async ingestion pipeline (Celery, Redis, Cloudflare R2) that fetches from the MediaWiki API and converts wikitext to markdown, scaling to **∼5M vectors per series** in pgvector — chosen over Pinecone/Weaviate to stay in one network hop.

•PromptForge — *Agentic LLM Orchestration & Cross-Model Benchmarking Platform* *Jan 2026 – Feb 2026*
Node.js, Express.js, BullMQ, Convex, OpenRouter, React.js

– Architected a creator–critic–evaluator agentic orchestration pipeline, scoring prompt clarity, constraint coverage, and token efficiency across versioned, regression-tracked revision trails — treating prompts as testable production artifacts.
– Implemented cross-model routing across OpenAI, Anthropic, Gemini, and open-source models via OpenRouter for objective quality, latency, and cost-per-output benchmarking.
– Built a fault-tolerant async backend using BullMQ workers with idempotent retries and Convex persistence, supporting reliable regression testing at scale without job loss on provider timeouts.
– Emitted structured prompt outputs with metadata — token budget, latency profile, expected output shape, version changelog — enabling prompts to be reviewed and deployed like code.

Technical Skills

Frontend: React.js, Next.js, Tailwind CSS, HTML, CSS
Backend: FastAPI, Node.js, Express.js, BullMQ, Celery, Redis, REST APIs, WebSockets
AI / ML: LLM Orchestration, Agentic Workflows, RAG Pipelines, Deploy-Gated Evaluation, pgvector, BM25, Reciprocal Rank Fusion, Multi-Agent Systems, Prompt Engineering, Behavioral Memory, TensorFlow, OpenRouter
Languages: Python, TypeScript, JavaScript, SQL, C++, Java
Databases: PostgreSQL (+ pgvector), MongoDB, Convex
Cloud & DevOps: AWS, Cloudflare R2, Git, GitHub, async job pipelines, CI-style eval gating
System Design: Multi-agent coordination, hybrid retrieval, event-driven architecture, snapshot-first serving